

Assignment 5: The Ocean

Practical information

Deadline: Sunday 13/3, 8pm

Resources:

- ERDA for file storage
- Jupyter for the Terminal to access DAG
- Nvidia profiler to determine the parallelisation bottlenecks
- Nvidia vGPUs for benchmarks on DAG

Handin:

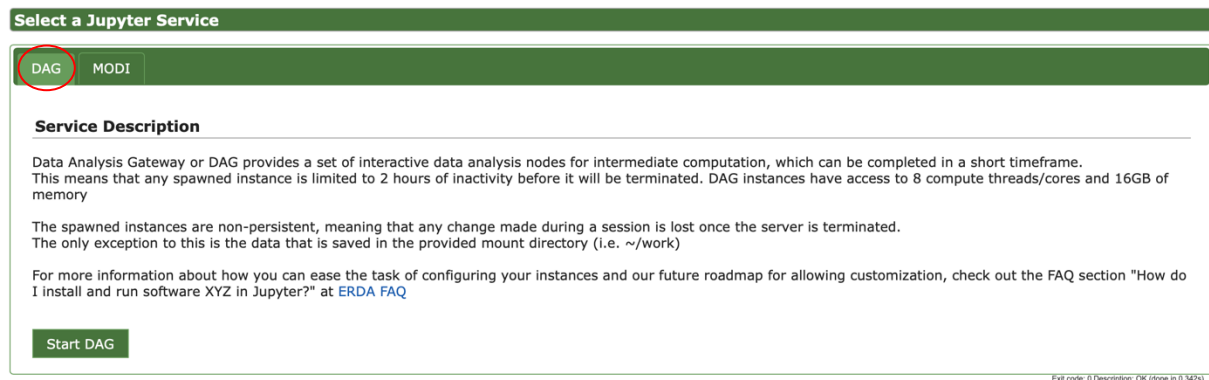
- Total assignment: a report of up to 3 pages in length (excluding the code)
- Use the template on Absalon to include your code in the report

Introduction

The Shallow Water (SW) model (section 13.3) is the simplest numerical representation of the ocean. Still, it has reasonable skills when used to predict the evolution of storm surges or Tsunamis. Moreover, it illustrates nicely the functioning and parallelization of stencil operations.

DAG

For this assignment we need `nvc++` to compile and DAG for source code profiling and running the benchmarks.



Select a Jupyter Service

DAG MODI

Service Description

Data Analysis Gateway or DAG provides a set of interactive data analysis nodes for intermediate computation, which can be completed in a short timeframe. This means that any spawned instance is limited to 2 hours of inactivity before it will be terminated. DAG instances have access to 8 compute threads/cores and 16GB of memory

The spawned instances are non-persistent, meaning that any change made during a session is lost once the server is terminated. The only exception to this is the data that is saved in the provided mount directory (i.e. `~/work`)

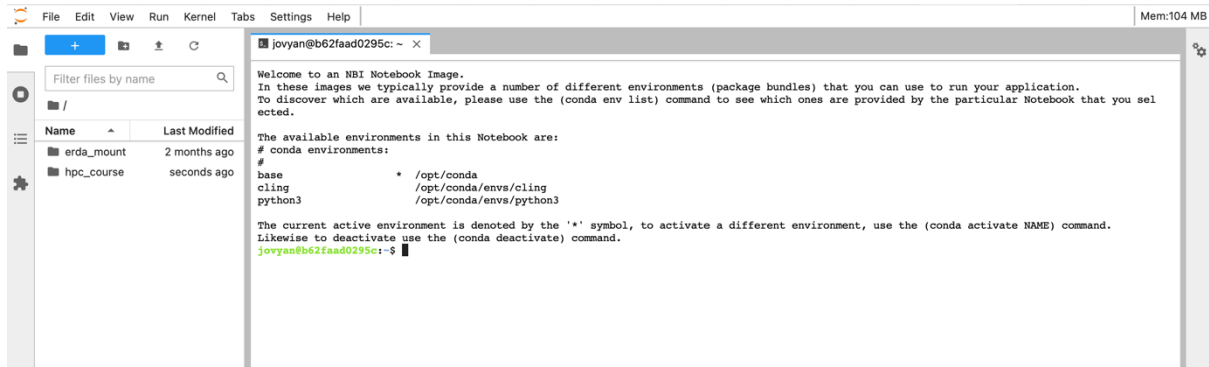
For more information about how you can ease the task of configuring your instances and our future roadmap for allowing customization, check out the FAQ section "How do I install and run software XYZ in Jupyter?" at [ERDA FAQ](#)

Start DAG

Exit code: 0 Description: OK (done in 0.342s)

You can read more about DAG in the user guide: <https://erda.dk/public/ucph-erda-user-guide.pdf>

Spin up a Jupyter session on DAG selecting the “HPC GPU notebook” notebook image. In the terminal (or the folder view on the right side) you can see a number of folders.



The different folders contain:

`erda_mount` : your own files.

`hpc_course` : course files.

Preparations

Start by copying the exercise to your storage area and enter in to the folder. You can write 'ls' to get a file listing of the folder.

```
cd erda_mount/HPPC
cp -a ~/hpc_course/module5 .
cd module5
ls
```

To be able to edit the files for the exercise navigate to the same folder in the file view. Here you can see 4 files:

```
Makefile
sw_parallel.cpp
sw_sequential.cpp
visualize.ipynb
```

Before you can run the code, you need to compile it. This can be done by running `make` in the terminal. The `sw_sequential.cpp` code is identical to `sw_parallel.cpp` there (with produced corresponding binaries `sw_sequential` and `parallel`) to give you a backup. The `visualize.ipynb` is for SW model output visualisation and analysis.

To run the code for the default 1000 time-steps on a grid of 512x512 on DAG and write the model output to a file, run

```
./sw_sequential
```

Nvidia profiler

NVIDIA profiler enables you to understand and optimize the performance of your OpenACC application. An example of command-line `nvprof` profiler output for parallelised SW model is given below. One can, for example, see a runtime of four compute kernels (`integrate_116_gpu`, `integrate_123_gpu`, `exchange_vertical_ghost_lines_100_gpu`, `exchange_horizontal_ghost_lines_87_gpu`) and time spent

for copying data from Host to Device and back, lines with [CUDA memcpy HtoD] and [CUDA memcpy DtoH], respectively.

```
jovyan@b62faad0295c:~/erda_mount/Teaching/hpc_course_private/module5$ nvprof ./sw_parallel --iter 500
==534== NVPROF is profiling process 534, command: ./sw_parallel --iter 500
checksum: 4117.75
elapsed time: 1.06693 sec
==534== Profiling application: ./sw_parallel --iter 500
==534== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 43.05% 16.155ms 500 32.310us 32.032us 33.183us integrate_116_gpu(Water&)
34.99% 13.129ms 500 26.258us 25.984us 26.880us integrate_123_gpu(Water&)
10.59% 3.9730ms 1000 3.9730us 3.9030us 14.720us exchange_vertical_ghost_lines_100_gpu(std::array<std::array<float,
unsigned long=512>, unsigned long=512>*)
10.02% 3.7595ms 1000 3.7590us 3.7110us 5.6640us exchange_horizontal_ghost_lines_87_gpu(std::array<std::array<float
, unsigned long=512>, unsigned long=512>*)
0.68% 254.56us 1 254.56us 254.56us 254.56us [CUDA memcpy HtoD]
0.67% 252.19us 1 252.19us 252.19us 252.19us [CUDA memcpy DtoH]
API calls: 69.07% 203.85ms 1 203.85ms 203.85ms 203.85ms cuDevicePrimaryCtxRetain
18.15% 53.572ms 7002 7.6510us 1.0040us 635.72us cuStreamSynchronize
7.98% 23.540ms 1 23.540ms 23.540ms 23.540ms cuMemHostAlloc
4.36% 12.860ms 3000 4.2860ms 3.4370us 533.67us cuLaunchKernel
0.28% 838.18us 1 838.18us 838.18us 838.18us cuMemAllocHost
0.08% 225.78us 2 112.89us 109.45us 116.33us cuMemAlloc
0.05% 148.92us 1 148.92us 148.92us 148.92us cuModuleLoadDataEx
0.01% 30.053us 1 30.053us 30.053us 30.053us cuMemcpyHtoDAsync
0.00% 12.201us 3 4.0670us 2.1660us 7.5830us cuEventRecord
0.00% 10.629us 1 10.629us 10.629us 10.629us cuMemcpyDtoHAsync
0.00% 8.4790us 2 4.2390us 1.8640us 6.6150us cuDeviceGetPCIBusId
0.00% 7.6780us 1 7.6780us 7.6780us 7.6780us cuPointerGetAttributes
0.00% 6.4870us 3 2.1620us 826ns 3.1350us cuEventCreate
0.00% 4.7920us 4 1.1980us 392ns 2.9810us cuModuleGetFunction
0.00% 4.2280us 10 422ns 156ns 2.1590us cuDeviceGetAttribute
0.00% 3.0310us 4 757ns 149ns 2.5130us cuDeviceGet
0.00% 2.6640us 3 888ns 211ns 1.8230us cuDeviceGetCount
0.00% 2.2560us 1 2.2560us 2.2560us 2.2560us cuEventSynchronize
0.00% 2.1050us 1 2.1050us 2.1050us 2.1050us cuCtxGetCurrent
0.00% 2.0000us 3 666ns 299ns 908ns cuCtxSetCurrent
0.00% 526ns 2 263ns 175ns 351ns cuDeviceComputeCapability
0.00% 231ns 1 231ns 231ns 231ns cuDriverGetVersion
OpenACC (excl): 19.91% 24.159ms 1 24.159ms 24.159ms 24.159ms acc_enter_data@sw_parallel.cpp:145
16.83% 20.430ms 1500 13.619us 1.7520us 57.933us acc_wait@sw_parallel.cpp:116
15.25% 18.508ms 1500 12.338us 1.7480us 637.24us acc_wait@sw_parallel.cpp:123
8.48% 10.290ms 2000 5.1440us 1.8420us 25.397us acc_wait@sw_parallel.cpp:100
8.45% 10.256ms 2000 5.1280us 1.8510us 25.614us acc_wait@sw_parallel.cpp:87
4.91% 5.9630ms 1000 5.9630us 4.6990us 535.39us acc_enqueue_launch@sw_parallel.cpp:87 (_Z38exchange_horizontal_gho
st_lines_87_gpuR5arrayIS_Iflm512EElm512EE)
4.37% 5.3077ms 1000 5.3070us 4.6310us 23.830us acc_enqueue_launch@sw_parallel.cpp:100 (_Z37exchange_vertical_ghos
t_lines_100_gpuR5arrayIS_Iflm512EElm512EE)
2.71% 3.2845ms 500 6.5680us 4.8360us 592.57us acc_enqueue_launch@sw_parallel.cpp:116 (_Z17integrate_116_gpu5Wat
```

Task 1: OpenACC parallelise the program (points 5)

The key challenge is to identify which parts of the code can reasonably be executed by the GPUs and to find suitable OpenACC directives and clauses for optimal parallelization. With the help of a profiler determine the bottlenecks. Play around a bit with the `#pragma` and see if you can improve on your first try. Thus, you need to save the profiler output of your various experiments. To get all 5 points, you should experiment around and attempt different paths for optimization using OpenACC.

Task 2: Strong and weak scaling (points 3)

Measure the weak and strong scaling of your programs. You should adjust `num_gangs()` to change how much of the GPU is actually used, and explain what this means for how the work actually maps to the physical hardware. Note that `num_gangs()` controls number of blocks NOT number of SMs.

For the weak scaling, you should change the grid size such that the calculations per thread stays constant, and for the strong scaling the grid size should be set such that the scaling from 1 to 2 is approximately linear like last week. Some key figures to note: You have 14 SMs available, the maximum number of threads per thread-block is 1024 and each multiprocessor can handle at most 2048 threads.

Task 3: Physics (points 2)

Test if the theoretically predicted phase speed of $c = (gH)^{1/2}$ is correctly reproduced by your model. Quantify and explain the possible differences between theory and simulation.

Bonus

The group who has the fastest parallelized code at 4pm with default settings for Nx, Ny and iter=10000 is rewarded a six-pack of beer.